Day 69 — AI Roadmap

Zecpath AI System — AI Roadmap

Current prototype baseline

• Modular resume parsing and normalization components.

• ATS scoring, screening, HR interview, technical interview and machine-test evaluation components.

• Role-weighted cross-round aggregation and explainable recommendation logic.

• Recruiter report formatting, in-process observability/audit utilities and synthetic end-to-end demo.

• Human review required; no autonomous hiring decision.

Roadmap

Phase 1 — Reliability and integration

• Define versioned contracts between each module.

• Replace in-memory API stores with a durable database and managed job queue.

• Add authentication, role-based authorization, secure upload validation, malware scanning and retention controls.

• Add integration tests against a deployed staging environment.

Phase 2 — Evaluation and evidence

• Create a consented, representative, privacy-reviewed evaluation dataset.

• Define task-specific metrics and baselines for parsing, matching and scoring.

• Measure calibration, error rates, subgroup performance and inter-rater agreement.

• Add reviewer override capture and auditability with access controls.

• Keep automated outputs advisory and provide contestability.

Phase 3 — Responsible AI and monitoring

• Conduct privacy, security, accessibility and legal reviews before pilot use.

• Monitor drift, missing data, latency, failure rates and score distributions.

• Add model/version lineage, alert routing and retention/deletion workflows.

• Validate accommodations and prevent protected attributes or proxies from improperly influencing ranking.

Phase 4 — Controlled pilot

• Pilot with trained recruiters in shadow mode before any workflow reliance.

• Compare system outputs with independently reviewed evidence.

• Gather stakeholder feedback and document acceptance criteria.

• Expand only after governance approval and measured benefit.

Proposed future capabilities (not implemented)

• Video interview analytics, emotion inference and candidate coaching are proposals only. Any future behavioral/video feature requires explicit consent, necessity assessment, bias validation, accessibility safeguards and legal review. Do not infer personality, honesty or job suitability from facial expression or gaze.

• Candidate-facing feedback and analytics dashboards require separate product and privacy design.

• Production-grade distributed observability, persistence and asynchronous processing remain future work.